


Class 7: Machine Learning

Izabelle Querubin

Example of K-means clustering

First step is to make up some data with a known structure, so we know what the answer should be.

```
tmp <- c(rnorm(30, mean = -3), rnorm(30, mean = 3))  
x <- cbind(x = tmp, y = rev(tmp))  
plot(x)
```



Now we have some structured data in `x`. Let's see if k-means is able to identify the two groups.

```
k <- kmeans(x, centers = 2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.968109	2.973144
2	2.973144	-2.968109

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 57.20792 57.20792
(between_SS / total_SS = 90.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Let's explore k:


```
k$size
```

[1] 30 30

k\$centers


	x	y
1	-2.968109	2.973144
2	2.973144	-2.968109

```
plot(x, col = k$cluster)
```




Now we can add the clusters centers:

```
plot(x, col = k$cluster)
points(k$centers, col = 'blue', pch = 15)
```



```
k_3 <- kmeans(x, centers = 3, nstart = 20)
plot(x, col = k_3$cluster)
```



Example of Hierarchical Clustering

Let's use the same data as before, which we stored in `x`. We will use the `hclust()` function.

```
clustering <- hclust(dist(x))  
clustering
```

Call:

```
hclust(d = dist(x))
```


Cluster method : complete

Distance : euclidean

Number of objects: 60


```
plot(clustering)
```

Cluster Dendrogram



Let's add a horizontal line:


Cluster Dendrogram



To get our results (i.e. membership vector), we need to “cut” the tree. The function for doing this is `cuttree()`.

[illegible]

Plotting this...



You can also “cut” your tree with the number of clusters you want:

```
cutree(clustering, k = 2)
```

[illegible]

Principal Component Analysis (PCA)

PCA of the UK Food

First was to read the data:

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267

Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this question?

Rows: 17; Columns: 4


To answer this question, you could use the `dim()` function.

Q2. Which approach to solving the 'row-names problem' mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Using `x <- read.csv(url, row.names=1)` is preferable and more robust since it is much simpler and more efficient.

Now we can generate some basic visualizations.

```
barplot(as.matrix(x), col = rainbow(nrow(x)))
```




Q3. Changing what optional argument in the `barplot()` function results in the following plot?

`beside = TRUE` —TRUE will render the plot below, FALSE will render the plot above.


Let's refine our barplot.

```
barplot(as.matrix(x), col = rainbow(nrow(x)), beside = TRUE)
```



Other visualizations that can be useful...

```
pairs(x, col = rainbow(nrow(x)), pch = 16)
```



Q5. Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

These plots are meant to compare population means and determine how significantly different they are from one another. A point that lies on the diagonal suggests that the value of that variable is similar in both populations.

Q6. What is the main difference between N. Ireland and the other countries of the UK in terms of this data-set?

Most of the data is clustered towards the bottom left of the plot, indicating that N. Ireland consumes different amounts of the different foods compared to the other countries of the UK.

Let's apply PCA (principal components analysis). For that, we need to use the command `prcomp()`. This function expects the transpose of our data.

```
#transpose_matrix <- t(s)
# pca <- prcomp(transpose_matrix)

pca <- prcomp(t(x))
```


```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's plot the PCA results!

```
plot(pca)
```



We need to access the results of the PCA analysis.

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"

$class
[1] "prcomp"
```

We can explore the `pca$x` dataframe:


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	2.532999	-105.768945	2.842865e-14
Wales	-240.52915	224.646925	56.475555	7.804382e-13
Scotland	-91.86934	-286.081786	44.415495	-9.614462e-13
N.Ireland	477.39164	58.901862	4.877895	1.448078e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


Plotting:

```
plot(x = pca$x[, 1], y = pca$x[, 2], xlab = "PC1", ylab = "PC2", xlim = c(-270,500))  
text(pca$x[, 1], pca$x[, 2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at the start of this document.

```
plot(x = pca$x[, 1], y = pca$x[, 2])
colors_countries <- c('orange', 'pink', 'blue', 'green')
text(x = pca$x[, 1], y = pca$x[, 2], colnames(x), col = colors_countries)
```




```
v <- round(pca$sdev^2/sum(pca$sdev^2) * 100)
v
```

```
[1] 67 29 4 0
```

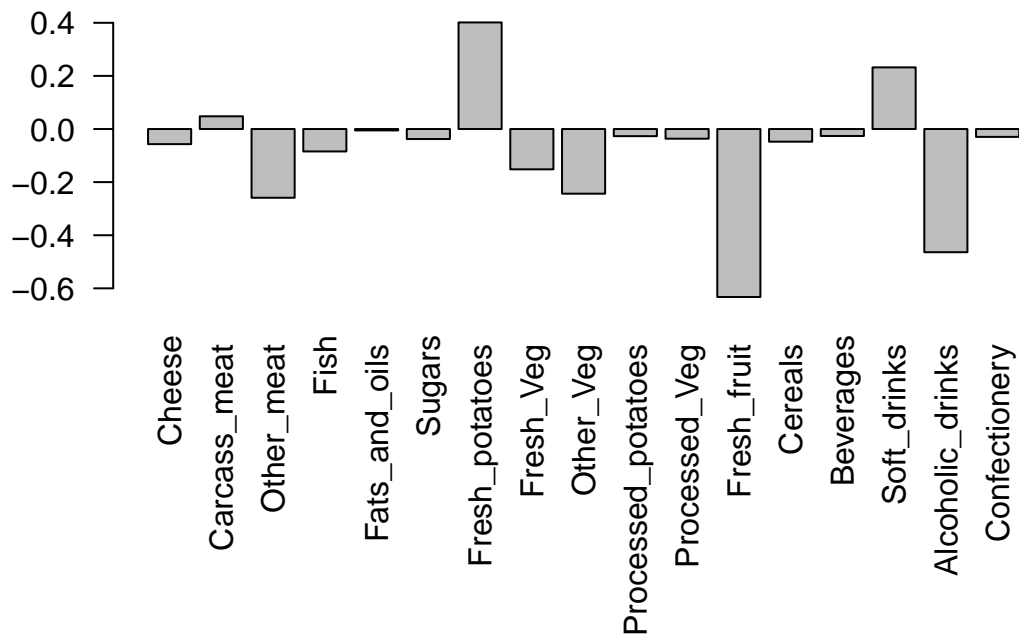
```
## or the second row here...
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	4.188568e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab = "Principal Component", ylab = "Percent Variation")
```




```
## Let's focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot(pca$rotation[,1], las = 2)
```




Q9. Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
## PC2 loadings plot
par(mar=c(10, 3, 0.35, 0))
barplot(pca$rotation[,2], las = 2)
```



The food with the largest positive loading score is mainly **Fresh_potatoes** and the food with the highest negative score is mainly **Soft_drinks**. PC2 tells us about the second principle component, which is the secondary trend or pattern that is orthogonal to PC1.

```
biplot(pca)
```

PCA of a RNA-Seq Dataset

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10. How many genes and samples are in this data set?

```
dim(rna.data)
```

```
[1] 100 10
```

I have 100 genes and 10 samples.

Let's apply PCA:

```
pca_rna = prcomp(t(rna.data), scale = TRUE)
summary(pca_rna)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251


	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.348e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

Let's plot the principle components 1 and 2.

```
plot(pca_rna$x[,1], pca_rna$x[,2], xlab = 'PC1', ylab = 'PC2')
```



```
plot(pca_rna, main = "Quick scree plot")
```




```
## Variance captured per PC
pca.var <- pca_rna$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
pca.var.per
```

```
[1] 92.6  2.3  1.1  1.1  0.8  0.7  0.6  0.4  0.4  0.0
```

```
barplot(pca.var.per, main = "Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab = "Principal Component", ylab = "Percent Variation")
```


Scree Plot



```
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca_rna$x[,1], pca_rna$x[,2], col = colvec, pch = 16,
     xlab = paste0("PC1 (", pca.var.per[1], "%)"),
     ylab = paste0("PC2 (", pca.var.per[2], "%)"))

text(pca_rna$x[,1], pca_rna$x[,2], labels = colnames(rna.data), pos = c(rep(4,5), rep(2,5)))
```




Using ggplot

```
library(ggplot2)


df <- as.data.frame(pca_rna$x)

# Our first basic plot
ggplot(df) +
  aes(PC1, PC2) +
  geom_point()
```



```
# Add a 'wt' and 'ko' "condition" column
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data), 1, 2)


p <- ggplot(df) +
  aes(PC1, PC2, label = samples, col = condition) +
  geom_label(show.legend = FALSE)
p
```



```
p + labs(title = "PCA of RNASeq Data",
  subtitle = "PC1 cleanly separates wild-tyoe from knock-out samples",
  x = paste0("PC1 (", pca.var.per[1], "%)"),
  y = paste0("PC2 (", pca.var.per[2], "%)"),
  caption = "Class example data") +
theme_bw()
```

PCA of RNASeq Data

PC1 cleanly separates wild-type from knock-out samples



Class example data